

Chapter 4

Trees

**South China University of
Technology**
College of Software Engineering

Huang Min


Chapter 4

Tree- part2

B-Tree

Beyond Binary Search Trees: Multi-Way Trees

- Example: B-tree of order 3 has 2 or 3 children per node



- Search for 8

B-Trees

B-Trees are **multi-way search trees** commonly used in database systems or other applications where data is stored externally on disks and keeping the tree shallow is important.

A **B-Tree of order M** has the following properties:


1. The **root** is either a leaf or has **between 2 and M children**.
2. All nonleaf nodes (except the root) have **between $\lceil M/2 \rceil$ and M children**.
3. **All leaves are at the same depth**.

All data records are stored at the leaves.
Internal nodes have “keys” guiding to the leaves.
Leaves store between $\lceil L/2 \rceil$ and L data records,
where L can be equal to M (default) or can be different.

B-Tree Details

Each (non-leaf) internal node of a B-tree has:


- › Between $\lceil M/2 \rceil$ and M children.
- › up to $M-1$ **keys** $k_1 < k_2 < \dots < k_{M-1}$



Keys are ordered so that:

$$k_1 < k_2 < \dots < k_{M-1}$$

Properties of B-Trees



Children of each internal node are "between" the items in that node.

Suppose subtree T_i is the i th child of the node:

all keys in T_i must be between keys k_{i-1} and k_i

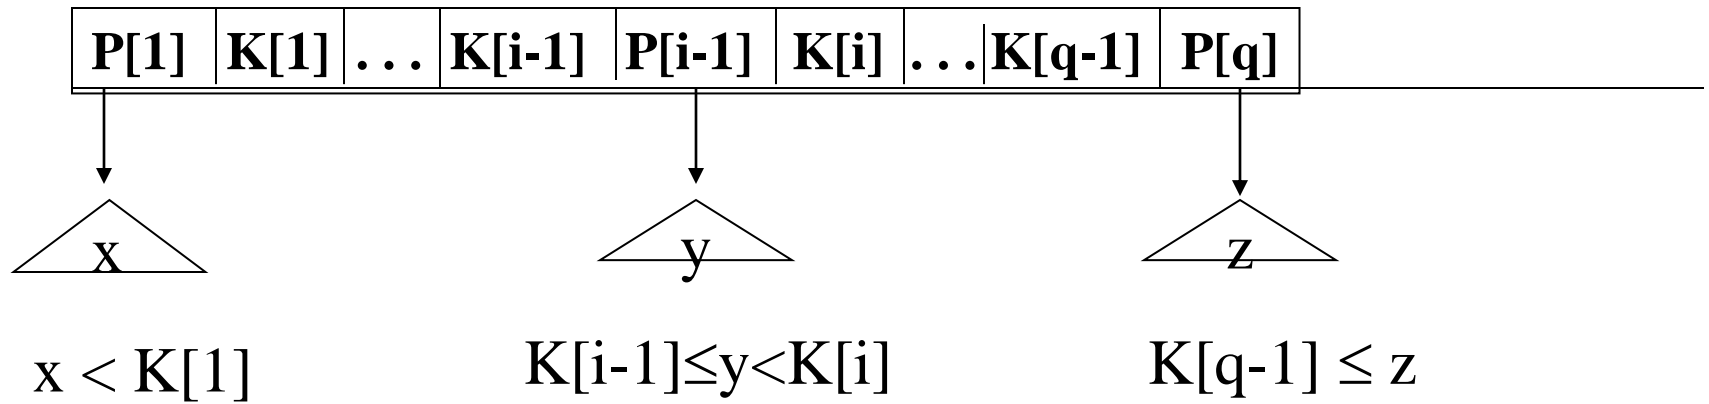
i.e. $k_{i-1} \leq T_i < k_i$

k_{i-1} is the smallest key in T_i


All keys in first subtree $T_1 < k_1$

All keys in last subtree $T_M \geq k_{M-1}$

B-Tree Nonleaf Node



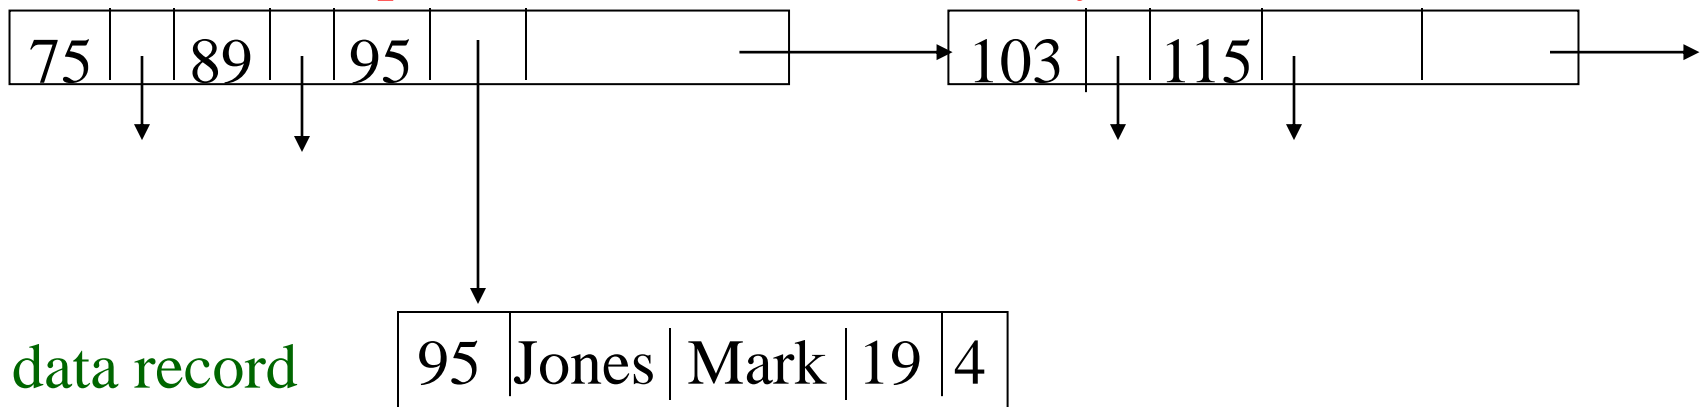
- The K s are keys
- The P s are pointers to subtrees.



Detailed Leaf Node Structure (B+ Tree)




- The Ks are keys (assume unique).
- The Rs are pointers to records with those keys.
- The Next link points to the next leaf in key order (**B+-tree**).



Searching in B-trees

- B-tree of order 3: also known as 2-3 tree (2 to 3 children)




- Examples: Search for 9, 14, 12
- Note: If leaf nodes are connected as a Linked List, B-tree is called a B+ tree – Allows sorted list to be accessed easily

Inserting into B-Trees


- Insert X: Do a Find on X and find appropriate leaf node
 - › If leaf node is not full, fill in empty slot with X
 - E.g. Insert 5
 - › If leaf node is full, **split** leaf node and adjust parents up to root node
 - E.g. Insert 9

Assume $M=L=3$,
so (6 7 8) is full.








Deleting From B-Trees

- Delete X : Do a find and remove from leaf
 - › Leaf underflows – borrow from a neighbor
 - E.g. 11
 - › Leaf underflows and can't borrow – merge nodes, delete parent
 - E.g. 17




Example of Insertions into a B+ tree with $M=3$, $L=2$


Insertion Sequence: 9, 5, 1, 7, 3, 12



B+-Tree Deletion




Example of a B⁺ -tree of order four.




Remove 12

Deletion from the B⁺ -tree via borrowing from a sibling.

B+-Tree Deletion




Example of a B⁺-tree of order four.



(a)

Remove 33



(b)

Deleting the record with key value 33 from the B⁺-tree via collapsing siblings.

AVL Trees


Binary Search Tree - Best Time

- All BST operations are $O(d)$, where d is tree depth
- minimum d is $d = \lfloor \log_2 N \rfloor$ for a binary tree with N nodes
 - › What is the best case tree?
 - › What is the worst case tree?
- So, best case running time of BST operations is $O(\log N)$

Binary Search Tree - Worst Time

- Worst case running time is $O(N)$
 - › What happens when you Insert elements in ascending order?
 - Insert: 2, 4, 6, 8, 10, 12 into an empty BST
 - › Problem: Lack of “balance”:
 - compare depths of left and right subtree
 - › Unbalanced degenerate tree

Balanced and unbalanced BST



Approaches to balancing trees

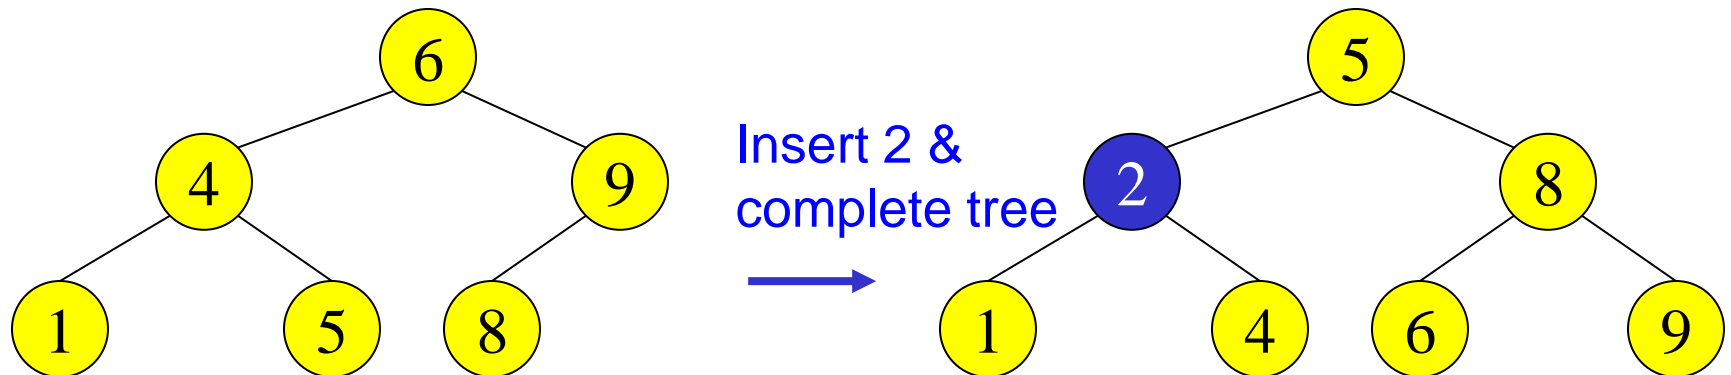
- Don't balance
 - › May end up with some nodes very deep
- Perfectly balance
 - › The tree must always be balanced perfectly
- Pretty good balance
 - › Only allow a little out of balance
- Adjust on access
 - › Self-adjusting

Balancing Binary Search Trees

- Many algorithms exist for keeping binary search trees balanced
 - › Adelson-Velskii and Landis (**AVL**) trees (height-balanced trees)
 - › **Splay trees** and other self-adjusting trees
 - › **B-trees** and other multiway search trees

Perfect Balance

- Want a **complete tree** after every operation
 - › tree is full except possibly in the lower right
- This is expensive
 - › For example, insert 2 in the tree on the left and then rebuild as a complete tree




AVL - Good but not Perfect Balance

- AVL trees are height-balanced binary search trees
- Balance factor of a node
 - › $\text{height}(\text{left subtree}) - \text{height}(\text{right subtree})$
- An AVL tree has balance factor calculated at every node
 - › For every node, heights of left and right subtree can differ by no more than 1
 - › Store current heights in each node

Height of an AVL Tree

- $N(h)$ = **minimum** number of nodes in an AVL tree of height h .
- **Basis**
 - › $N(0) = 1, N(1) = 2$
- **Induction**
 - › $N(h) = N(h-1) + N(h-2) + 1$
- **Solution** (recall Fibonacci analysis)
 - › $N(h) \geq \phi^h$ ($\phi \approx 1.62$)




Height of an AVL Tree


- $N(h) \geq \phi^h$ ($\phi \approx 1.62$)
- Suppose we have n nodes in an AVL tree of height h .
 - › $n \geq N(h)$ (because $N(h)$ was the minimum)
 - › $n \geq \phi^h$ hence $\log_{\phi} n \geq h$ (relatively well balanced tree!!)
 - › $h \leq 1.44 \log_2 n$ (i.e., Find takes $O(\log n)$)

Node Heights

Tree A (AVL)



Tree B (AVL)




height of node = h

balance factor = $h_{\text{left}} - h_{\text{right}}$


empty height = -1

Node Heights after Insert 7

Tree A (AVL)



Tree B (not AVL)





balance factor
 $1 - (-1) = 2$

height of node = h
balance factor = $h_{\text{left}} - h_{\text{right}}$
empty height = -1

Insert and Rotation in AVL Trees

- Insert operation may cause balance factor to become 2 or -2 for some node
 - › only nodes on the path from insertion point to root node have possibly changed in height
 - › So after the Insert, go back up to the root node by node, updating heights
 - › If a new balance factor (the difference $h_{\text{left}} - h_{\text{right}}$) is 2 or -2, adjust tree by *rotation* around the node

Single Rotation in an AVL Tree



Insertions in AVL Trees

Let the node that needs rebalancing be α .

There are 4 cases:

Outside Cases (require single rotation) :

1. Insertion into **left** subtree **of left** child of α .
2. Insertion into **right** subtree **of right** child of α .


Inside Cases (require double rotation) :

3. Insertion into **right** subtree **of left** child of α .
4. Insertion into **left** subtree **of right** child of α .


The rebalancing is performed through four separate rotation algorithms.

AVL Insertion: Outside Case


Consider a valid
AVL subtree




AVL Insertion: Outside Case




AVL Insertion: Outside Case



Single right rotation



Outside Case Completed




“Right rotation” done!
 (“Left rotation” is mirror
 symmetric)

AVL property has been restored!


AVL Insertion: Inside Case

Consider a valid
AVL subtree




AVL Insertion: Inside Case

Inserting into Y
destroys the
AVL property
at node j



Does “right rotation”
restore balance?


AVL Insertion: Inside Case



“Right rotation”
does not restore
balance... now k is
out of balance


AVL Insertion: Inside Case

Consider the structure
of subtree Y...




AVL Insertion: Inside Case


Y = node i and
subtrees V and W




AVL Insertion: Inside Case



Double rotation : first rotation




Double rotation : second rotation




Double rotation : second rotation

right rotation complete

Balance has been restored



Implementation




No need to keep the height; just the difference in height, i.e. the **balance** factor; this has to be modified on the path of insertion even if you don't perform rotations

Once you have performed a rotation (single or double) you won't need to go back up the tree

Single Rotation

```
RotateFromRight(n : reference node pointer) {  
  p : node pointer;  
  p := n.right;  
  n.right := p.left;  
  p.left := n;  
  n := p  
}
```


You also need to modify
the heights or balance
factors of n and p



Double Rotation

- Implement Double Rotation in two lines.

```
DoubleRotateFromRight(n : reference node pointer)
{
  ???
}
```



Insertion in AVL Trees

- Insert at the leaf (as for all BST)
 - › only nodes on the path from insertion point to root node have possibly changed in height
 - › So after the Insert, go back up to the root node by node, updating heights
 - › If a new balance factor (the difference $h_{\text{left}} - h_{\text{right}}$) is 2 or -2 , adjust tree by *rotation* around the node


Insert in BST

```
Insert(T : reference tree pointer, x : element) : integer
{
  if T = null then
    T := new tree; T.data := x; return 1; //the links to
                                         //children are null
  case
    T.data = x : return 0; //Duplicate do nothing
    T.data > x : return Insert(T.left, x);
    T.data < x : return Insert(T.right, x);
  endcase
}
```

Insert in AVL trees


```
Insert(T : reference tree pointer, x : element) : {  
  if T = null then  
    {T := new tree; T.data := x; height := 0; return;}  
  case  
    T.data = x : return ; //Duplicate do nothing  
    T.data > x : Insert(T.left, x);  
                  if ((height(T.left) - height(T.right)) = 2){  
                    if (T.left.data > x ) then //outside case  
                      T = RotatefromLeft (T);  
                    else //inside case  
                      T = DoubleRotatefromLeft (T);}  
    T.data < x : Insert(T.right, x);  
                  code similar to the left case  
  Endcase  
  T.height := max(height(T.left), height(T.right)) + 1;  
  return;  
}
```

Example of Insertions in an AVL Tree

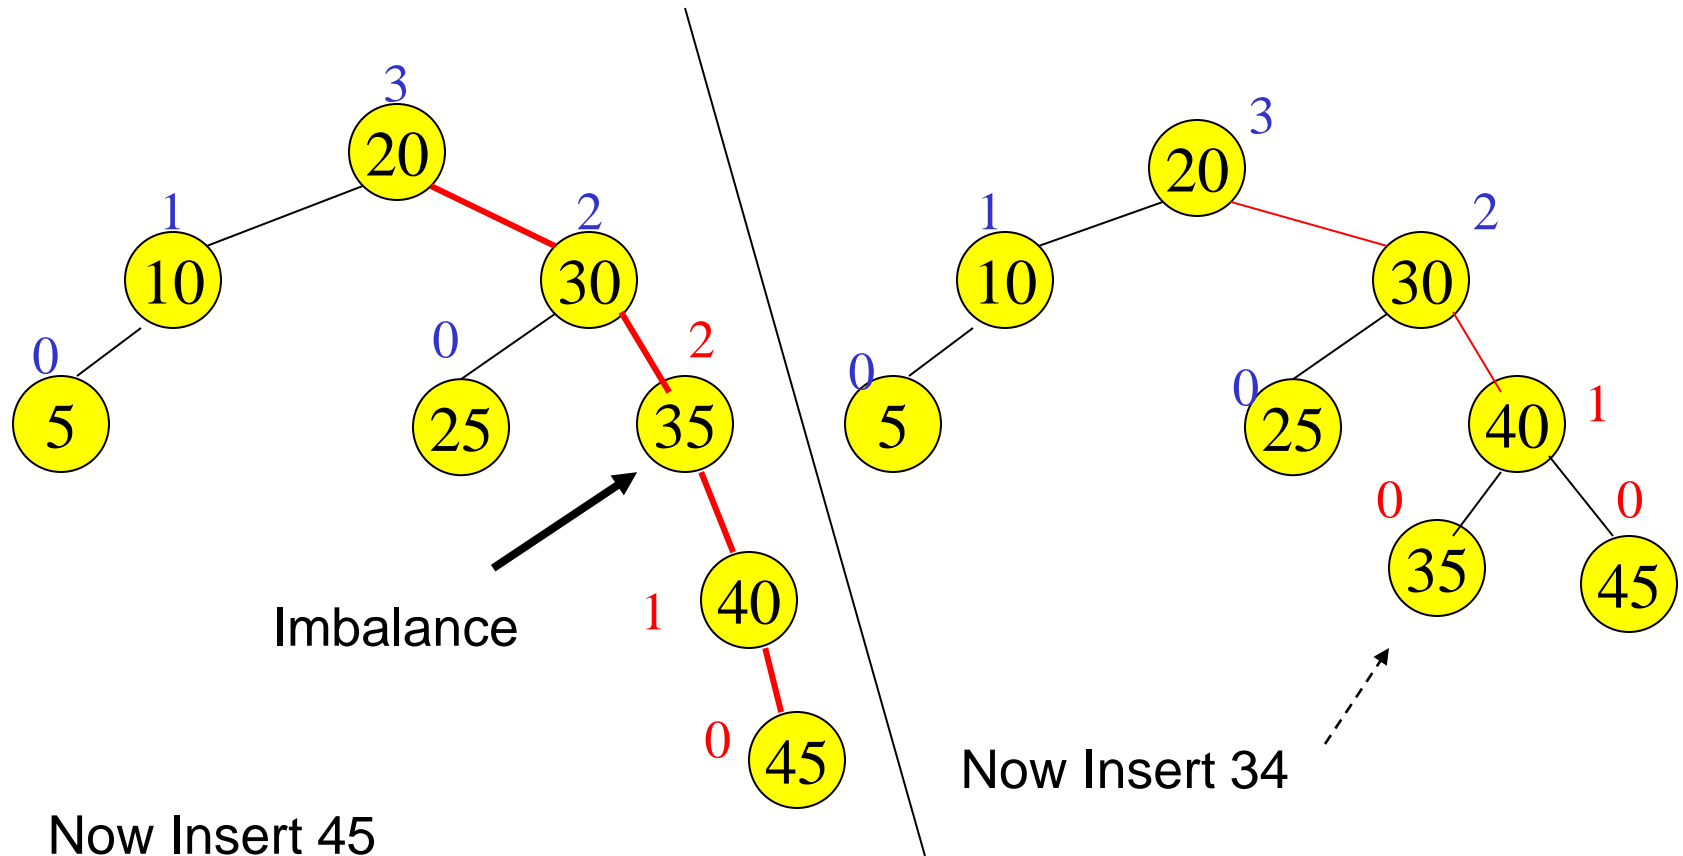


Insert 5


Example of Insertions in an AVL Tree



Single rotation (outside case)



Double rotation (inside case)



AVL Tree Deletion

- Similar but more complex than insertion
 - › Rotations and double rotations needed to rebalance
 - › Imbalance may propagate upward so that many rotations may be needed.

Splay Trees

Self adjusting Trees





- Ordinary binary search trees have no balance conditions
 - › what you get from insertion order is it
- Balanced trees like AVL trees enforce a balance condition when nodes change
 - › tree is always balanced after an insert or delete
- Self-adjusting trees get reorganized over time as nodes are accessed
 - › Tree adjusts after insert, delete, or find

Splay Trees

- Splay trees are tree structures that:
 - › Are not perfectly balanced all the time
 - › Data most recently accessed is near the root. (principle of locality; 80-20 “rule”)
- The procedure:
 - › After node X is accessed, perform “splaying” operations to bring X to the root of the tree.
 - › Do this in a way that leaves the tree more balanced as a whole

Splay Tree Terminology


- Let X be a non-root node with ≥ 2 ancestors.
 - P is its parent node.
 - G is its grandparent node.




Zig-Zig and Zig-Zag

Parent and grandparent
in same direction.

zig-zig




Parent and grandparent
in different directions.



Splay Tree Operations


1. Helpful if nodes contain a **parent** pointer.



2. When X is accessed, apply one of **six** rotation routines.
 - Single Rotations (X has a P (the root) but no G)
ZigFromLeft, ZigFromRight
 - Double Rotations (X has both a P and a G)
ZigZigFromLeft, ZigZigFromRight
ZigZagFromLeft, ZigZagFromRight

Zig at depth 1 (root)


- “Zig” is just a **single rotation**, as in an AVL tree
- Let R be the node that was accessed (e.g. using Find)



- ZigFromLeft moves R to the top → faster access next time

Zig at depth 1


- Suppose Q is now accessed using Find



- `ZigFromRight` moves `Q` back to the top


Zig-Zag operation

- “Zig-Zag” consists of **two rotations of the opposite direction** (assume R is the node that was accessed)




Zig-Zig operation


- “Zig-Zig” consists of two single rotations of the same direction (R is the node that was accessed)




Decreasing depth - "autobalance"




(a)



(b)



(c)



(d)

Find(T)



Find(R)




Splay Tree Insert and Delete



- Insert x
 - › Insert x as normal then splay x to root.
- Delete x
 - › Splay x to root and remove it. (note: the node does not have to be a leaf or single child node like in BST delete.) Two trees remain, right subtree and left subtree.
 - › Splay the max in the left subtree to the root
 - › Attach the right subtree to the new root of the left subtree.

Example Insert


- Inserting in order 1,2,3,...,8
- Without self-adjustment



With Self-Adjustment

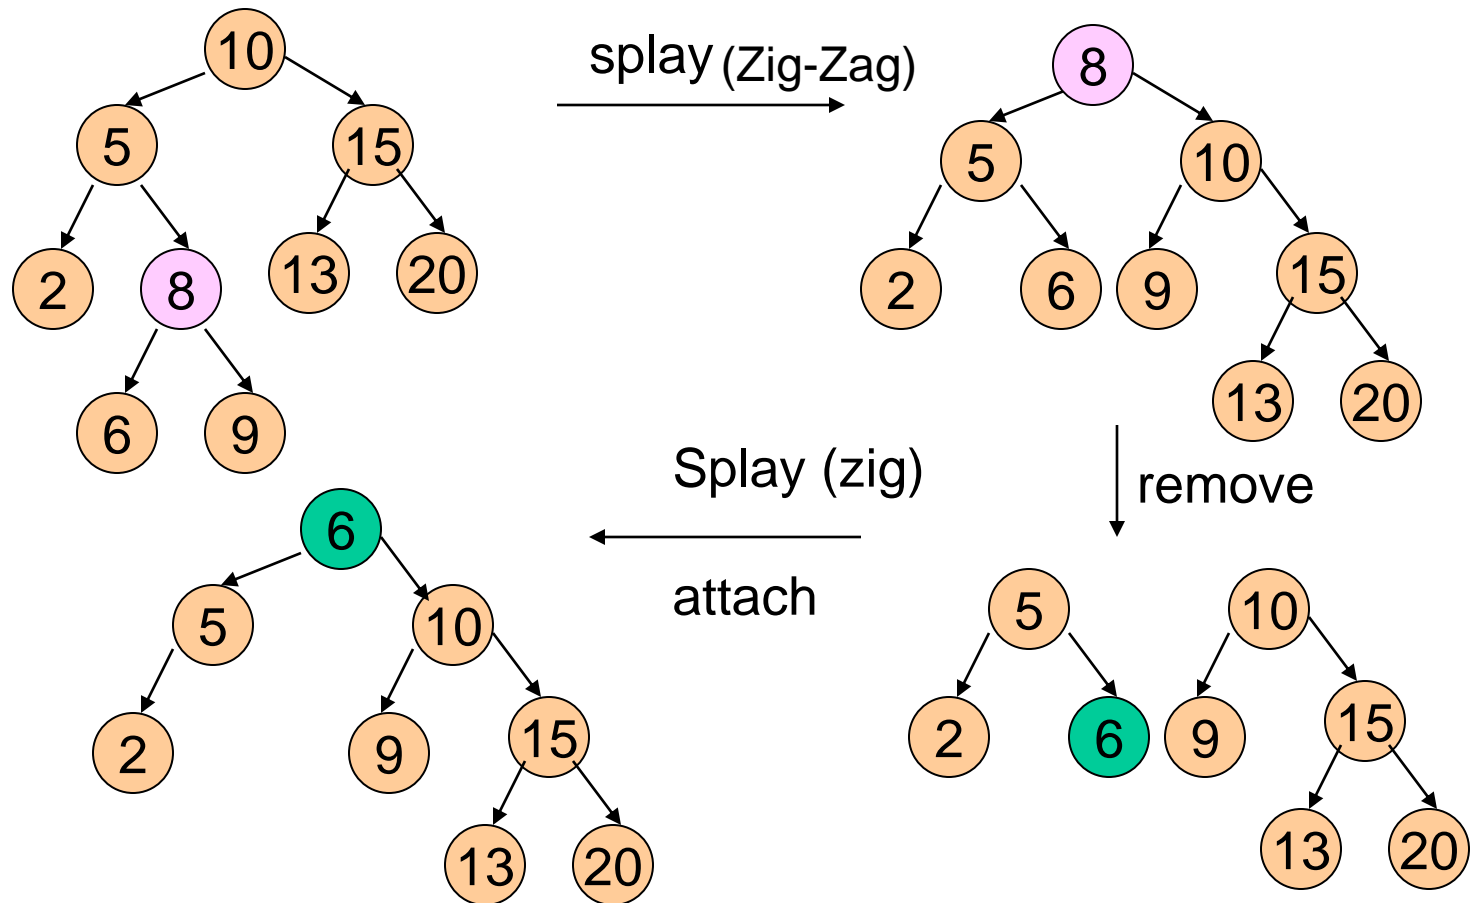


With Self-Adjustment



Each Insert takes $O(1)$ time therefore $O(n)$ time for n Insert!!

Example Deletion



Analysis of Splay Trees

- Splay trees tend to be balanced
 - › M operations takes time $O(M \log N)$ for $M \geq N$ operations on N items. (proof is difficult)
 - › Amortized $O(\log n)$ time.
- Splay trees have good “locality” properties
 - › Recently accessed items are near the root of the tree.
 - › Items near an accessed one are pulled toward the root.

Summary of Search Trees

- Problem with Binary Search Trees: Must keep tree balanced to allow fast access to stored items
- **AVL trees:** Insert/Delete operations keep tree balanced
- **Splay trees:** Repeated Find operations produce balanced trees
- **Multi-way search trees** (e.g. B-Trees):
 - › More than two children per node allows shallow trees; all leaves are at the same depth.
 - › Keeping tree balanced at all times.
 - › Excellent for indexes in database systems.

Summary

1. 二叉树的定义


定义： 是 n ($n \geq 0$) 个结点的有限集合，由一个根结点以及两棵互不相交的、分别称为**左子树**和**右子树**的二叉树组成。

逻辑结构： 一对二 (1: 2)


基本特征：

- ① 每个结点最多只有两棵子树 (不存在度大于2的结点) ;
- ② 左子树和右子树次序不能颠倒 (有序树) 。

基本形态：



具有3个结点的二叉树可能有几种不同形态?



具有n个结点的二叉树可能有几种不同形态?

答：设具有n个结点的所有不同形态的二叉树有 $b[n]$ 种，则 $b[n] = C(2n, n) / (n + 1)$ ($n=1, 2, 3, \dots$)

证明如下：

考虑n个结点。除去根，剩下n-1个结点.对左子树有 $b[k]$ 种方式。对右子树有 $b[n-1-k]$ 种方式，由乘法原理，则 $b[n] = \sum_{k=0}^{n-1} (b[k] * b[n-1-k])$ 由于 $b[1]=1$ 而这正好是Catalan数。

Catalan数（卡特兰数）：

令 $h(0)=1, h(1)=1$, catalan数满足递归式：

$h(n) = h(0)*h(n-1) + h(1)*h(n-2) + \dots + h(n-1)h(0)$ (其中 $n \geq 2$)

该递推关系的解为：

$$h(n) = C(2n, n) / (n + 1) \quad (n=1, 2, 3, \dots)$$

2. 二叉树的性质

性质1: 在二叉树的第 i 层上至多有 2^i 个结点 ($i \geq 0$) 。

性质2: 高度为 k 的二叉树至多有 $2^k - 1$ 个结点 ($k > 0$) 。

性质3: 具有 k 个节点的完全二叉树的高度为 $\lceil \log_2(k+1) \rceil$ ($k \geq 0$) 。

问：高度为9的二叉树中至少有_____个结点。

A) 2^9

B) 2^8

C) 9


D) $2^9 - 1$

答案：C

3. 二叉树的存储结构

一、顺序存储结构

按二叉树的结点“自上而下、从左至右”编号，用一组连续的存储单元存储。



[0]	A
[1]	B
[2]	C
[3]	D
[4]	E
[5]	F
[6]	G
[7]	H
[8]	I

问：顺序存储后能否复原成唯一对应的二叉树形状？

答：若是**完全二叉树**则可以做到**唯一复原**。



而且有规律：下标值为 i 的双亲，其左孩子的下标值必为 $2i+1$ ，其右孩子的下标值必为 $2i+2$ （即性质5）

例如，对应[2]的两个孩子必为[5]和[6]，即C的左孩子必是F，右孩子必为G。

不是完全二叉树怎么办?

答：一律转为完全二叉树！

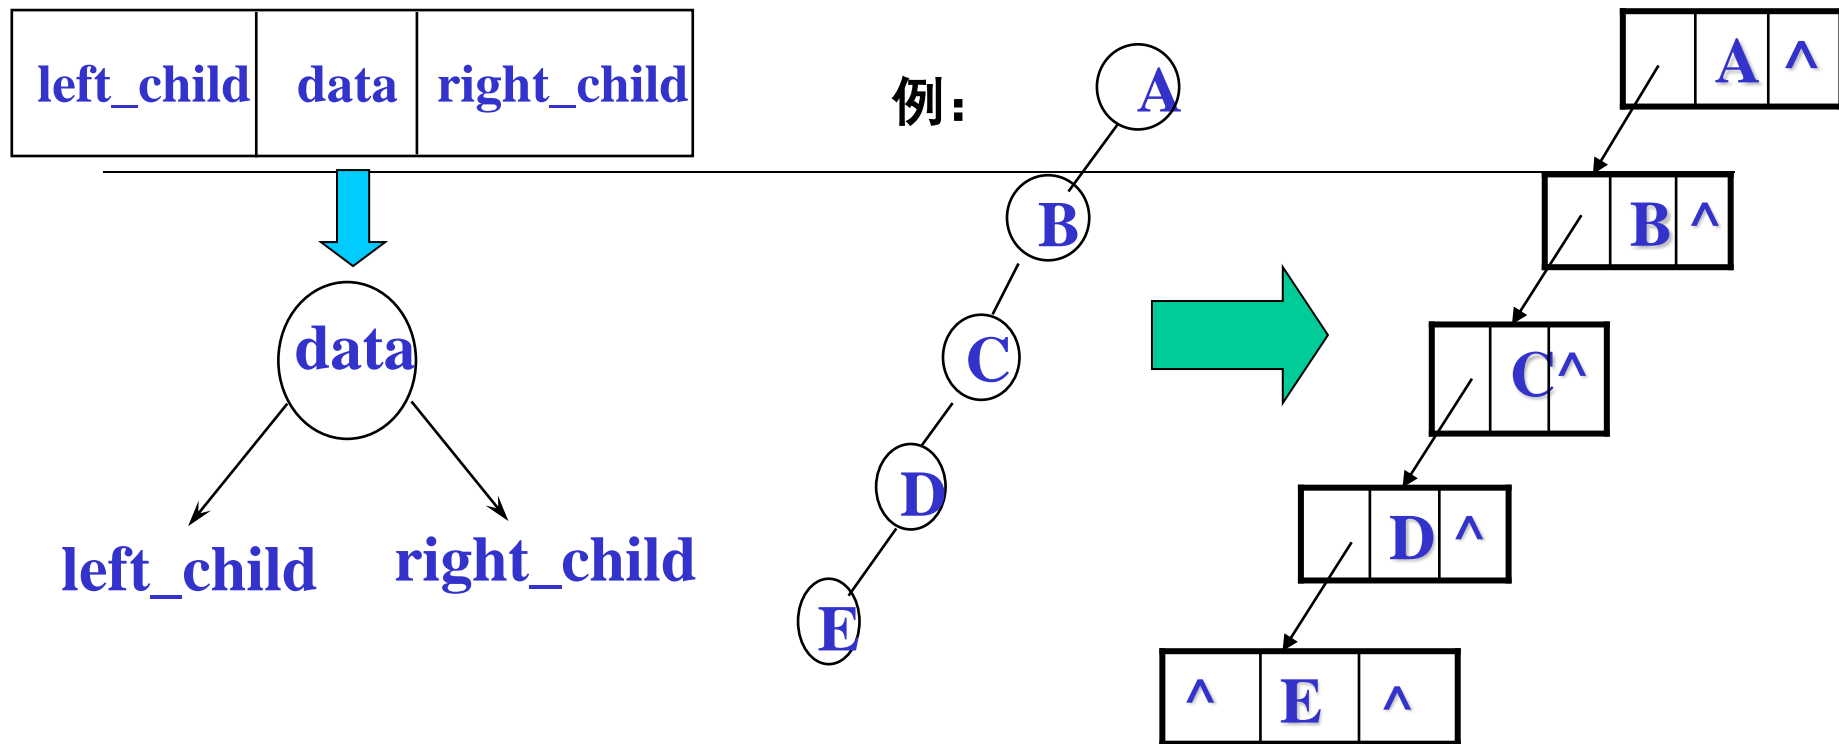
将各层空缺处统统补上“虚结点”，其内容为空。



[0]	A
[1]	B
[2]	^
[3]	C
[4]	^
[5]	^
[6]	^
[7]	D
[8]	^
.	...
[15]	E

缺点：①浪费空间；②插入、删除不便

二、链式存储结构 用二叉链表即可方便表示。



一般从根结点开始存储。

(相应地，访问树中结点时也只能从根开始)

注：如果需要倒查某结点的双亲，可以再增加一个双亲域（直接前趋）指针，将二叉链表变成三叉链表。

4、遍历二叉树 (Traversing Binary Tree)

遍历定义——指按某条搜索路线遍访每个结点且不重复（又称周游）。

遍历用途——它是树结构插入、删除、修改、查找和排序运算的前提，

是二叉树一切运算的基础和核心。

遍历规则——

❖ 二叉树由根、左子树、右子树构成，定义为**D、L、R**

❖ D、L、R的组合定义了六种可能的遍历方案：

LDR, LRD, DLR, DRL, RDL, RLD

❖ 若限定先左后右，则有三种实现方案：

DLR


先(根)序遍历

LDR

中(根)序遍历

LRD

后(根)序遍历



先序遍历的结果是: **A B D E C**

中序遍历的结果是: **D B E A C**

后序遍历的结果是: **D E B C A**

遍历的算法实现：用递归形式

5. 二叉查找树

特点：“左小右大”，按中序遍历得到由小到大的排列

主要操作

检索(find)——

折半查找

插入(insert) ——

通过折半查找找到插入的位置（插入某个叶结点或在待插入方向上没有子结点的分支结点）

删除最小值结点(deletemin)——

删除整个树中最左边的结点，若该结点有右子树，则将其父结点中原来指向被删结点的指针改为指向其右子树。

删除给定值结点(remove)——若被删除节点的左右子结点非空，则用其右子树中的最小节点取代被删除结点。

优点：提高检索、插入、删除等操作的效率，平均情况下 $\theta(\log n)$. 82