Computer Organization & Architecture Chapter 8 – Virtual Memory

Zhang Yang 张杨 cszyang@scut.edu.cn Autumn 2021

Contents of this lecture

- 8.8 Virtual Memory
 - What is Virtual Memory?
 - Motivation for Virtual Memory
 - □ Terminology of Virtual Memory
 - □ Implementation of Virtual Memory
 - □ Paging
 - □ Summary

What is Virtual Memory?

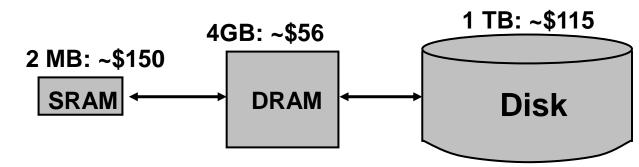
- If you think it's there, and it's there...it's real.
- If you think it's not there, and it's there...it's transparent.
- If you think it's there, and it's not there...it's imaginary.
- If you think it's not there, and it's not there...it's nonexistent.
- Virtual memory is imaginary memory: it gives you the illusion of a memory arrangement that's not physically there.

Motivations for Virtual Memory

- Use physical DRAM as cache for the disk.
 - Address space of a process can exceed physical memory size.
 - □ Sum of address spaces of multiple processes can exceed physical memory (more common modern case).
- Simplify memory management.
 - □ Multiple processes resident in main memory.
 - Each with its own address space.
 - □ Only "active" code and data is actually in memory.
 - Allocate more memory to process as needed.
- Provide protection.
 - One process can't interfere with another.
 - Because they operate in different address spaces.
 - □ User process cannot access privileged information.
 - Different sections of address spaces have different permissions.

Motivation#1: DRAM as cache for the disk

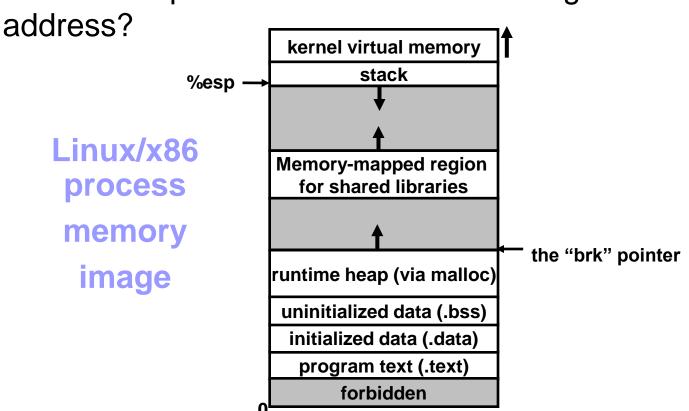
- Full address space is quite large:
 - □ 32-bit addresses: ~4,000,000,000 (4 billion) bytes
 - □ 64-bit addresses: ~16,000,000,000,000,000,000 (16 quintillion) bytes
- Disk storage is ~250X cheaper than DRAM storage.
 - □ 1 TB of DRAM: ~ \$28,000 (667 MHz, late 2008 prices)
 - □ 1 TB of disk: ~ \$115
- To get cost-effective access to large amounts of data, bulk of data must be stored on disk.



Motivation#2: Memory Management (1)

- Motivation#2: Memory Management
 - Multiple processes can reside in physical memory.
 - □ How do we resolve address conflicts?

What if two processes access something at same

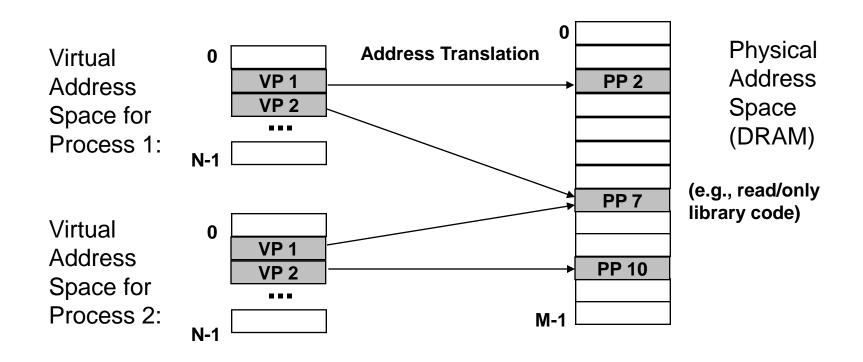


Motivation#2: Memory Management (2)

- Motivation#2: Memory Management (ctd.)
 - □ Solution: Separate Virtual Address Spaces
 - Virtual and physical address spaces divided into equal-sized blocks.
 - Blocks are called "pages" (both virtual and physical).
 - Each process has its own virtual address space.
 - Operating system controls how virtual pages are assigned to physical memory.

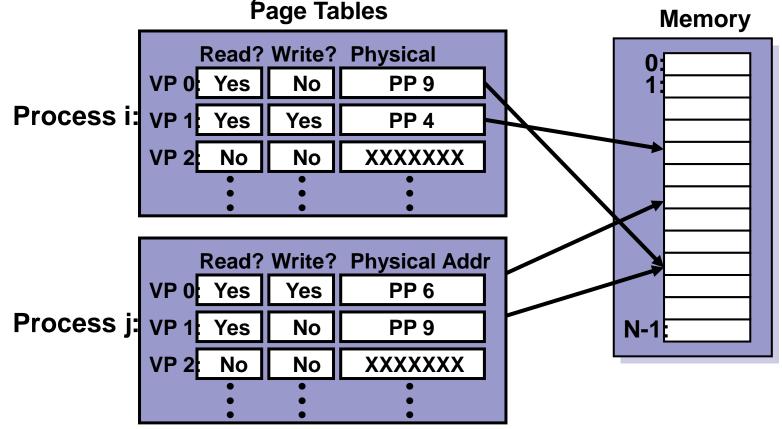
Motivation#2: Memory Management (3)

- Motivation#2: Memory Management (ctd.)
 - □ Solution: Separate Virtual Address Spaces (ctd.)



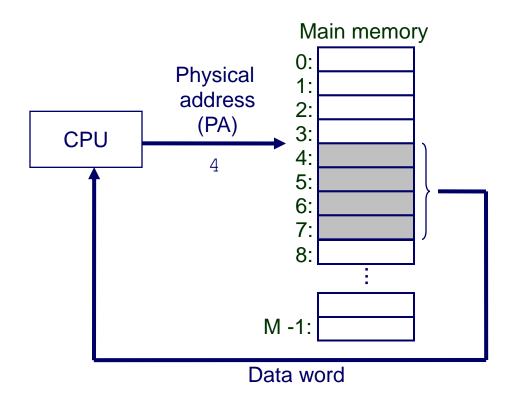
Motivation#3: Protection

- Page table entry contains access-rights information.
 - □ Hardware enforces this protection (trap into OS if violation occurs).



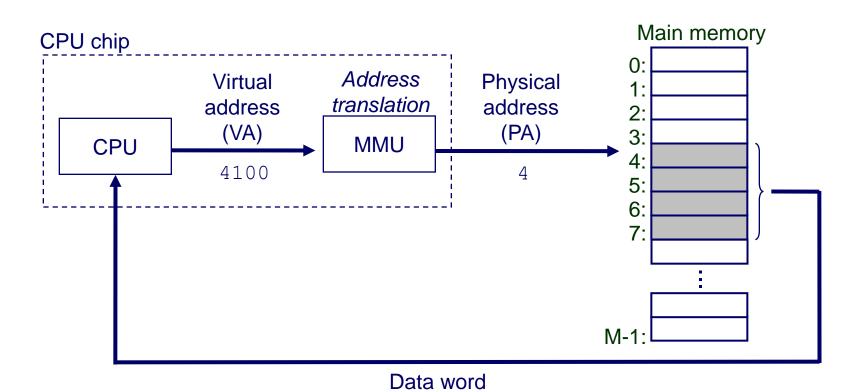
A System Using Physical Addressing

 Used by many digital signal processors and embedded microcontrollers in devices like phones and PDAs.



A System Using Virtual Addressing

 Used by all modern desktop and laptop microprocessors.

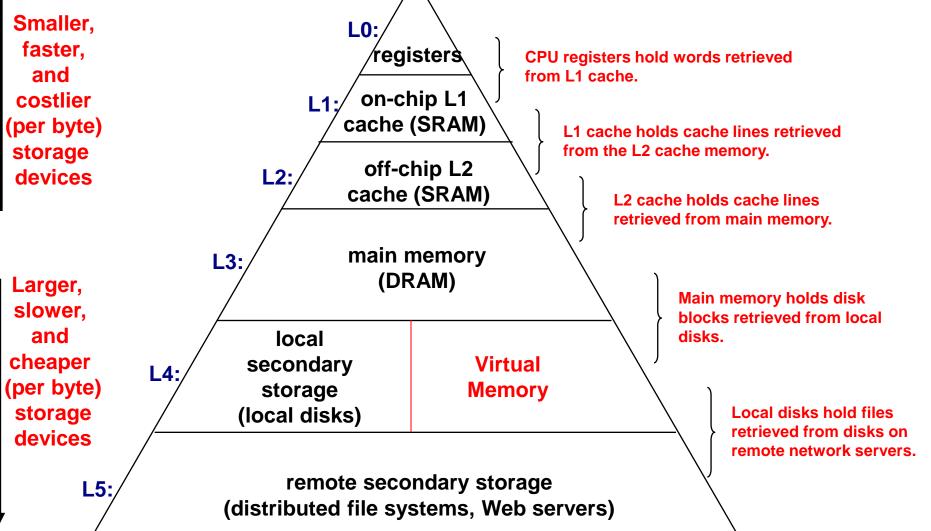


Terminology of VM

- Physical memory memory as seen by the processor
- Logical memory memory as seen by the process
- Physical Address real location in physical memory; identifies actual storage
- Virtual Address conventional addressing used by a program which the OS must translate into a physical address

Recall: Memory Hierarchy

A Figure of Memory Hierarchy

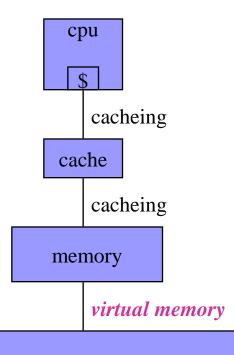


Why Use Disks?

- Multiprogramming
 - More than one application running at one time.
 - May need more than available main memory to store all needed programs and data.
 - May want to share data between applications.
- Cheap, large (but, slow)

VM Definition

- Not a physical device but an abstract concept.
- It's just another level in the cache/memory hierarchies.
- Virtual memory is the name of the technique that allows us to view main memory as a cache of a larger memory space (on disk).
 - Automatic storage allocation



disk

Typical Organization of VM

- MMU
 - Memory Management Unit
- OS
 - Control MMU
- Address Translation
 - The hardware converts
 virtual addresses into
 physical addresses via an
 OS-managed lookup table
 (page table)

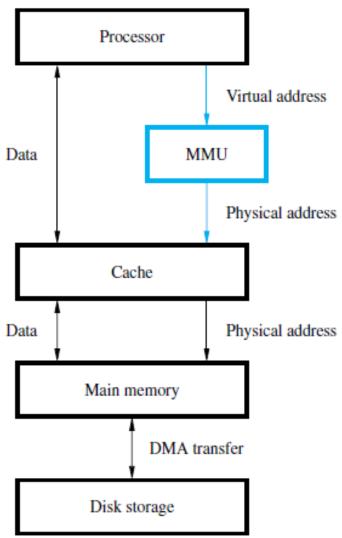


Figure 8.24 Virtual memory organization.

Benefits of VM

- Large address space: Virtual address to Physical Address mapping.
- Simplify relocation: the same program can run in any location in the physical memory.
- Reduce time to start a program: not all code and data need to be in the physical memory before a program can start.
- Relieve programmers from burden of overlays: VM automatically manages the two levels of memory hierarchy represented by main memory and secondary storage.

Implementation of VM

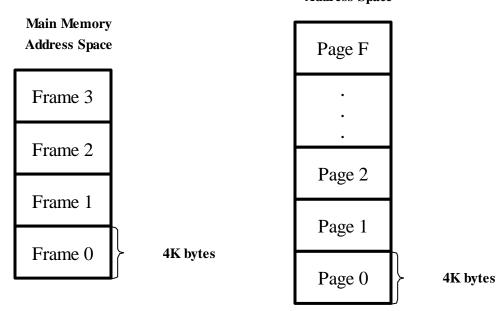
- Three Implementations
 - Paging
 - Segmentation
 - Segmentation with Paging

Paging

- Address Translation
- Page Replacement
- Write Policy
- Internal Fragmentation
- Page Size

Paging (1)

Organization of Virtual Memory and Main Memory
 Memory
 Address Space



- Page: a block of consecutive words
- Page Frame: an area in the main memory that can hold one page

Paging (2)

- Organization of Virtual Memory and Main Memory (ctd.)
 - □ Note
 - Pages constitute the basic unit of information that is moved between the main memory and the disk.
 - Pages commonly range from 2K to 16K bytes in length.
 - Pages should not be too small or too large.

Address Structure

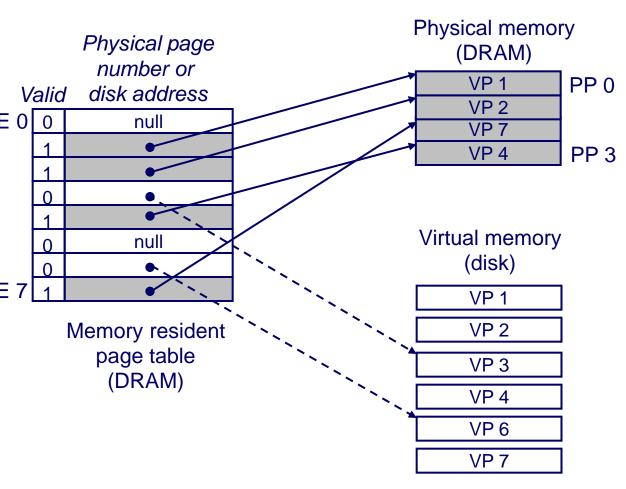
- Virtual Address (2 fields)
 - Virtual page number (high-order bits)
 - □ Offset (low-order bits)
 - Specify the location of a particular byte(or word) within a page.
 - Physical Address (2 fields)
 - Physical page number (high-order bits)
 - Offset

Page Table

Page Table

A page table
is an array of
page table PTE 0 0
entries
(PTEs) that
maps virtual
pages to
physical
pages.

Kernel data structure in DRAM



Page Table Entry

- Access Control Bit
 - □ Read, Write, Execute, etc
- Present Bit (Valid Bit)
 - present in main memory (or not)
- Dirty Bit
 - If the page has been modified
- Use Bit Access Present Dirty Use Page Number
 - □ Recently used?
- Page Number
 - physical page number OR pointer to secondary storage

Address Translation (1)

Translate Virtual Memory Address to Physical Address
Virtual address from processor

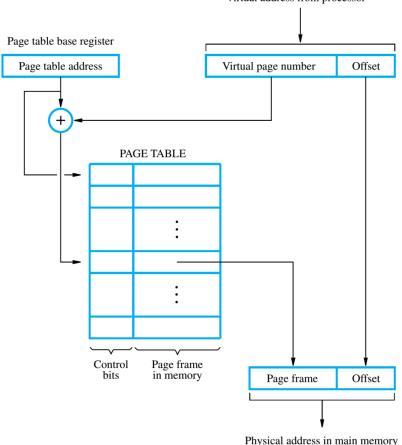
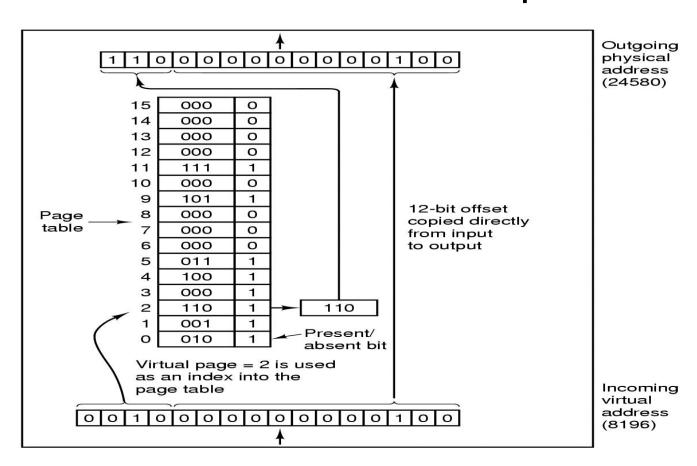


Figure 8.25 Virtual-memory address translation.

Address Translation (2)

Address Translation Example



Page Hit and Page Fault (1)

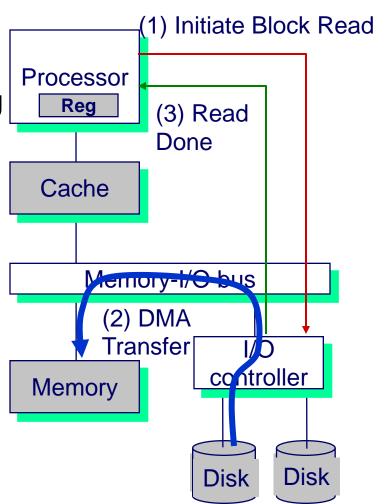
- Page Hit
 - A page hit is a reference to a VM word that is in physical (main) memory.
- Page Fault (Similar to "Cache Miss")
 - A page fault is caused by a reference to a VM word that is not in physical (main) memory.
 - Example: A instruction references a word contained in VP 3, a miss that triggers a page fault exception.
 - What if an object is on disk rather than in memory?
 - Page table indicates that the virtual address is not in memory.
 - OS trap handler is invoked, moving data from disk into memory.
 - Current process suspends, others can resume.
 - OS has full control over placement, etc.

Page Hit and Page Fault (2)

- Servicing a Page Fault
 - Make space in memory by writing physical page to disk.
 - Page Frames
 - Replacement policy?
 - Load page
 - Loading pages could waste processor time, use DMA.
 - DMA allows processor to do something else.
 - OS updates the process's page table.
 - Desired data is in memory for process to resume.



- Servicing a Page Fault (ctd.)
 - Processor signals disk controller.
 - Read block of length P starting at disk address X and store starting at memory address Y.
 - Read occurs
 - Direct Memory Access (DMA)
 - Under control of I/O controller
 - Controller signals completion
 - Interrupts processor.
 - OS resumes suspended process.



TLB (1)

- Disadvantage of Page Table
 - Every memory access by a program can take at least twice as long: one memory access to obtain the physical address and a second access to get the data.

TLB

- A small cache which is incorporated in MMU consists of a small portion of the page table.
- The contents of the TLB must be coherent with the contents of page tables in the main memory.

TLB (2)

AddressTranslationwith TLB

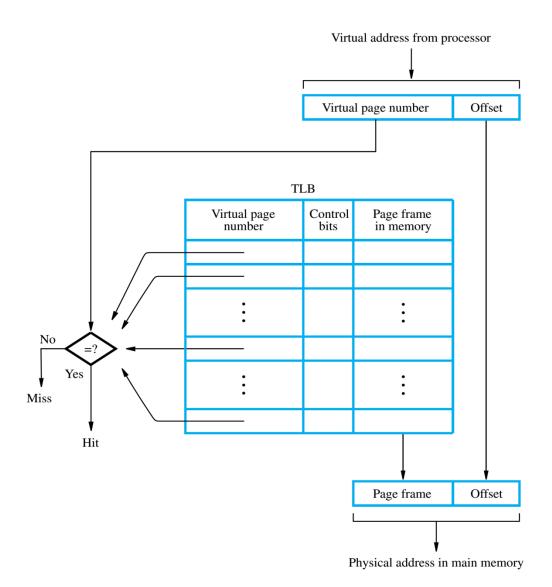
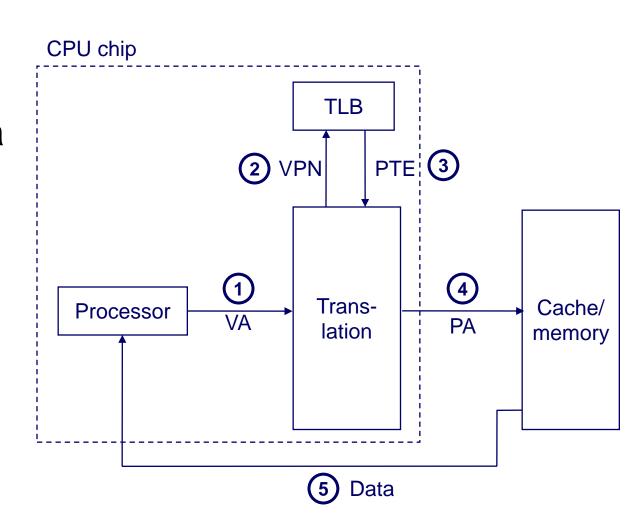


Figure 8.26 Use of an associative-mapped TLB.

TLB (3)

TLB Hit

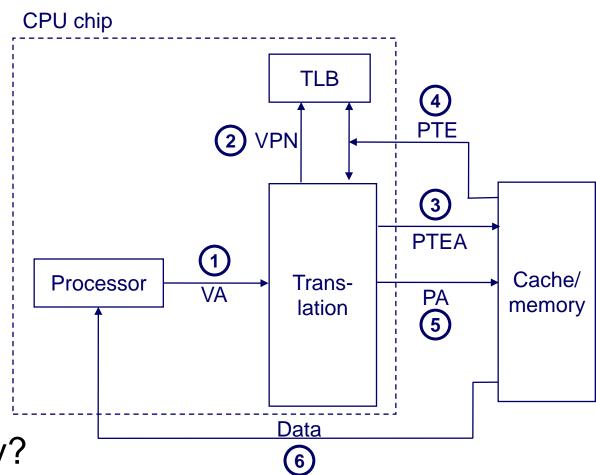
 A TLB hit eliminates a memory access.



TLB (4)

TLB Miss

- A TLB miss incurs an additional memory access (the PTE).
- □ Fortunately,TLB missesare rare. Why?

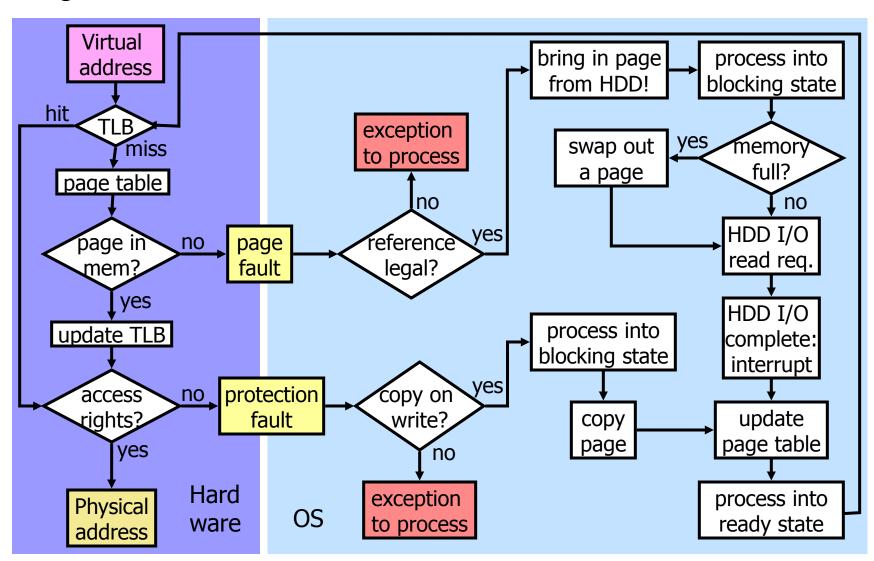


TLB (5)

- Design Issues
 - Accessed Frequently
 - Speed is important.
 - Minimize TLB Miss
 - Can be fully-associative, set-associative, or direct mapped.
 - Often fully associative, can be set associative.
 - Sized to maximize hits, but make timing.
 - Usually not more than 128, 256 entries (associativity).

Address Translation Summary

Figure



Page Replacement

- LRU (needs use bits)
- FIFO
- Software (OS) implementation

Write Policy

- Write back
- Write through is not suitable for virtual memory. Why?

Internal Fragmentation

- Disadvantage of Paging: Internal Fragmentation
 - Last page is unlikely to be full.
 - □ Example: If the program and data need 26000 bytes on a machine with 4096 bytes per page, the first 6 page will be full, totaling 6×4096=24576 bytes, and the last page will contain 26000—24576=1424 bytes. So 2672 bytes will be wasted.
 - Conclusion
 - If the page size is n bytes, the average amount of space wasted in the last page of a program by internal fragmentation will be n/2 bytes.

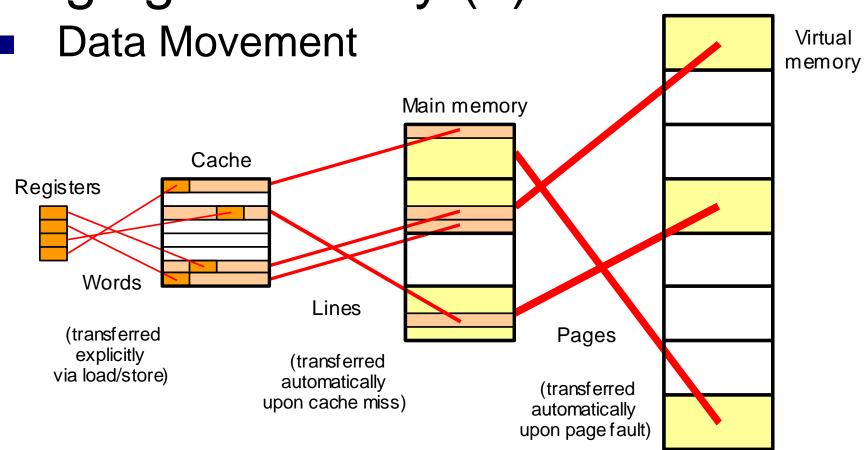
Page Size (1)

- Choosing a page size is a question of balancing forces that favor a larger page size versus those favoring a smaller size.
 - Advantages of Choosing Large Page Size
 - The size of the page table is inversely proportional to the page size: Memory can therefore be saved by making the pages bigger
 - Transferring larger pages to or from secondary storage, possibly over a network, is more efficient than transferring smaller pages.

Page Size (2)

- Advantage of Choosing Smaller Page Size
 - □ The main reason for choosing smaller page is conserving storage. A small page size will result in less wasted storage when a contiguous region of virtual memory is not equal in size to a multiple of the page size.





Data movement in a memory hierarchy.

Paging Summary (2)

Comparing the 2 levels of hierarchy

<u>Cache</u>	Virtual Memory
Block or Line	Page
Miss	Page Fault
Block Size: 32-64B	Page Size: 4K-16KB
Placement: Direct Mapped, N-way Set Associative	Fully Associative
Replacement: LRU or Random	Least Recently Used (LRU) approximation
Write Thru or Back	Write Back
How Managed: Hardware	Hardware + Software (Operating System)

Summary (1)

- 知识点: Virtual Memory
 - □ 了解Motivations of VM
 - □掌握What is VM
 - 了解Implementation of VM: Paging,
 Segmentation, Segmentation with paging
- 知识点: Paging
 - □掌握Address Translation
 - □理解Page hit and Page fault
 - □理解TLB的用途

Summary (2)

- 知识点: Paging (ctd.)
 - □理解Page table and page table entry
 - □ 了解Page replacement
 - □理解Write policy
 - □理解Internal fragmentation
 - □掌握Page size
 - Smaller vs. bigger

M

Exercise (1)

■ 简答: What are advantages and disadvantages of choosing smaller page size in a paging system?

Solution

Advantages: A small page size will result in less wasted storage when a contiguous region of virtual memory is not equal in size to a multiple of the page size.

□ Disadvantages:

- (1) The size of the page table is inversely proportional to the page size. Memory can therefore be wasted by making the pages smaller.
- (2) Transferring smaller pages to or from secondary storage, possibly over a network, is less efficient than transferring larger pages.



Exercise (2)

■ 简答: Under what circumstances do page faults occur? Describe the actions taken by the operating system when a page fault occurs.

Solution

- A page fault occurs when an access to a page that has not been brought into main memory takes place.
- □ The operating system verifies the memory access, aborting the program if it is invalid. If it is valid, a free frame is located and I/O is requested to read the needed page into the free frame. Upon completion of I/O, the process table and page table are updated and the instruction is restarted.



Exercise (3)

- 计算: Assume the following memory system:
 - □ Virtual addresses are 14 bits
 - □ Physical addresses are 12 bits
 - ☐ The page size is 64 bytes
 - The first 16 entries in the page table are as follows: (VPN means virtual page number, PPN means physical page number)

_
_

	PPN	Valid
80	13	1
09	17	1
0A	09	1
0B	_	0
0C	1E	1
0D	2D	1
0E	11	1
0F	0D	1



Exercise (4)

- (1) What is the format of a virtual address? In other words, how many bits are used for the virtual page number and how many bits are used for the virtual page offset?
- (2) What is the format of a physical address? In other words, how many bits are used for the physical page number and how many bits are used for the physical page offset within a page?
- (3) Show how this memory system translates the virtual address 0x030D into the corresponding physical address. Write the translation process!



Exercise (5)

Solution:

- □ (1)Virtual page offset: 6 bits, virtual page number: 8 bits
- □ (2)Physical page offset: 6 bits, physical page number: 6 bits
- □ (3)0x030D= (00 0011 0000 1101)₂, so the virtual page number is $(00001100)_2$ =0x0C. Use this VPN to lookup the page table, get the PPN=1E.

So the physical address = $(011110001101)_2=0x078D$

м

Homework (1)

- 补充题
- 4. Consider the memory system with the following specifications:

Byte-addressable.

Virtual address space: 4G bytes.

Main memory size: 16M bytes.

Cache size: 512K bytes.

Page size: 64K bytes.

Block size: 32 bytes

Mapping Strategies: Main Memory to Cache: 8-way set associative; Hard Disk to Main Memory: fully associative.

Virtual address is first translated to physical address. Then, it accesses the cache memory using the physical address.

м

Homework (2)

- 补充题(续)
- (1) How many sets are there in the cache memory?
- (2) How long is the tag field of the cache?
- (3) Given a virtual address 0547AF33 (hexadecimal), its corresponding virtual page is stored in physical page 3B (hexadecimal).
- What is its physical address under such mapping?
- 2 Which set can this address be possibly found in the cache?
- 3 Which byte does this address point to out of the 32 bytes in a block?